

Online and Reinforcement Learning (2025)

Home Assignment 7

Davide Marchi 777881

Contents

1	Short Questions	2
2	Offline Evaluation of Bandit Algorithms - the Practical Part	2
3	Grid-World: Continual and undiscounted	2
4	An Empirical Evaluation of UCB Q-learning	4

1 Short Questions

1. **True.** *Justification:* In a finite average-reward MDP with a finite diameter the MDP is communicating (in fact, ergodic if every state is reachable from every other under some policy). This implies that the long-run average reward (gain) is independent of the starting state.
2. **False.** *Justification:* Finite diameter guarantees that there exists a policy which can reach any state from any other in a finite expected number of steps. However, this does not mean that every arbitrary choice of actions will eventually reach every state.
3. **False.** *Justification:* In an ergodic MDP the optimal gain is unique (i.e., independent of the initial state), but the optimal bias function is determined only up to an additive constant. Hence, the bias is not uniquely defined in an absolute sense.
4. **False.** *Justification:* A PAC-MDP algorithm guarantees that the number of ε -bad (i.e., non- ε -optimal) steps is bounded with high probability. However, it does not imply that there exists a finite time after which every subsequent policy is ε -optimal; occasional exploration may still yield suboptimal actions.

2 Offline Evaluation of Bandit Algorithms - the Practical Part

3 Grid-World: Continual and undiscounted

In this question, we model the 4-room grid-world as an average-reward MDP and solve it using Value Iteration (VI).

(i) Implementation of Value Iteration for Average-Reward MDPs

The following Python function implements VI for average-reward MDPs. The algorithm computes an approximation of the optimal gain g^* , the bias function b^* (normalized such that $\min_s b^*(s) = 0$), and the optimal policy.

```
def average_reward_vi(mdp, epsilon=1e-6, max_iter=10000):  
  
    # Initialize the value function arbitrarily (here, zeros)  
    V = np.zeros(mdp.nS)  
  
    for iteration in range(max_iter):  
        V_next = np.zeros(mdp.nS)  
        # Update V_next for each state s using the Bellman operator:  
        # V_next(s) = max_a [ R(s, a) + sum_x P(x|s,a) * V(x) ]  
        for s in range(mdp.nS):  
            action_values = []
```

```

        for a in range(mdp.nA):
            q_sa = mdp.R[s, a] + np.dot(mdp.P[s, a], V)
            action_values.append(q_sa)
        V_next[s] = max(action_values)

    # Compute the difference (increment) vector
    diff = V_next - V
    # The span (max difference minus min difference) is our stopping
criterion
    span_diff = np.max(diff) - np.min(diff)
    V = V_next.copy()

    if span_diff < epsilon:
        print(f"Converged after {iteration+1} iterations with span {
span_diff:.2e}.")
        break
    else:
        print("Warning: Maximum iterations reached without full
convergence.")

    # Estimate the optimal gain g* as the average of the maximum and
minimum differences.
    gain = 0.5 * (np.max(diff) + np.min(diff))
    # The bias function is approximated by normalizing V (bias is defined
up to an additive constant)
    bias = V - np.min(V)

    # Derive the optimal policy: for each state, choose the action
maximizing:
    # Q(s, a) = R(s, a) + sum_x P(x|s,a) * V(x)
    policy = np.zeros(mdp.nS, dtype=int)
    for s in range(mdp.nS):
        best_val = -np.inf
        best_a = 0
        for a in range(mdp.nA):
            q_sa = mdp.R[s, a] + np.dot(mdp.P[s, a], V)
            if q_sa > best_val:
                best_val = q_sa
                best_a = a
        policy[s] = best_a

    return policy, gain, bias

```

Listing 1: Average-Reward Value Iteration Function

(ii) Visualization of the Optimal Policy

Running the above function on the grid-world produced the following output:

- **Optimal Gain:** $g^* = 0.07563$

- **Span of the Optimal Bias:** $\text{sp}(b^*) = 0.92437$

The optimal policy for the grid-world is represented by the following table. The arrows indicate the action chosen in each state (with walls marked as “Wall”):

Wall	Wall	Wall	Wall	Wall	Wall	Wall
Wall	→	→	→	↓	↓	Wall
Wall	↓	↑	Wall	↓	←	Wall
Wall	↓	Wall	Wall	↓	Wall	Wall
Wall	→	→	→	→	↓	Wall
Wall	→	↑	Wall	→	↑	Wall
Wall	Wall	Wall	Wall	Wall	Wall	Wall

Table 1: Optimal Policy Grid: Arrows indicate the optimal actions.

(iii) Interpretation of $1/g^*$

In this grid-world, the optimal gain g^* represents the long-run average reward per time step. Hence, the quantity

$$\frac{1}{g^*} \approx \frac{1}{0.07563} \approx 13.22$$

can be interpreted as the average number of steps needed to collect one unit of reward under the optimal policy. In other words, the agent gathers one reward approximately every 13 time steps, reflecting the efficiency of the optimal strategy in this continual setting.

4 An Empirical Evaluation of UCB Q-learning

In this exercise we implement UCB Q-learning in the 5-state RiverSwim environment. The algorithm is run with the parameters:

$$\gamma = 0.92, \quad \varepsilon = 0.13, \quad \delta = 0.05, \quad T = 2 \times 10^6.$$

The starting state is sampled uniformly and we set

$$H = \left\lceil \frac{1}{1-\gamma} \log \frac{1}{\varepsilon} \right\rceil, \quad b(k) = \sqrt{\frac{H}{k} \log \left(SA \log(k+1)/\delta \right)}.$$

The performance is measured by the cumulative number of ε -bad steps,

$$n(t) = \sum_{\tau=1}^t \mathbf{1}\{V^{\pi_\tau}(s_\tau) < V^*(s_\tau) - \varepsilon\},$$

where $V^*(s)$ is computed via value iteration. In our experiments the value iteration procedure yields:

$$V^* = [2.83439639, 3.45056954, 4.27771502, 5.31104628, 6.594788].$$

Below is the complete implementation of the UCB Q-learning routine.

```
def run_ucb_ql(T, gamma, epsilon_eval, delta, H, V_star, record_interval=
1000, seed=None):
    if seed is not None:
        np.random.seed(seed)
    env = riverswim(5)
    # Sample starting state uniformly.
    env.s = np.random.randint(0, env.nS)
    s = env.s
    S = env.nS
    A = env.nA
    R_max = 1.0
    optimistic_init = R_max / (1.0 - gamma)
    Q = np.full((S, A), optimistic_init)
    Q_hat = np.copy(Q)
    N = np.ones((S, A))
    n_bad = 0
    times = []
    n_values = []
    # Cache for policy evaluation
    policy_value_cache = {}
    for t in range(1, T + 1):
        # Select action greedily w.r.t. Q_hat.
        a = np.argmax(Q_hat[s, :])
        new_s, r = env.step(a)
        k = N[s, a]
        alpha_k = (H + 1) / (H + k)
        bonus = np.sqrt((H / k) * np.log(S * A * np.log(k + 1) / delta))
        # Q-update with exploration bonus.
        Q[s, a] = (1 - alpha_k) * Q[s, a] + alpha_k * (r + bonus + gamma
* np.max(Q_hat[new_s, :]))
        Q_hat[s, a] = min(Q_hat[s, a], Q[s, a])
        N[s, a] += 1
        # Evaluate current greedy policy.
        policy = np.argmax(Q_hat, axis=1)
        policy_tuple = tuple(policy)
        if policy_tuple not in policy_value_cache:
            V_policy = evaluate_policy(policy, env, gamma)
            policy_value_cache[policy_tuple] = V_policy
        else:
            V_policy = policy_value_cache[policy_tuple]
        if V_policy[new_s] < V_star[new_s] - epsilon_eval:
            n_bad += 1
        s = new_s
        if t % record_interval == 0:
            times.append(t)
            n_values.append(n_bad)
    return np.array(times), np.array(n_values)
```

Listing 2: Complete UCB Q-learning routine

(i) Sample Path of $n(t)$

A single run of the algorithm produces a sample path of the cumulative ε -bad steps $n(t)$. Figure 1 shows the sample path. This plot illustrates the steady accumulation of bad steps over time.

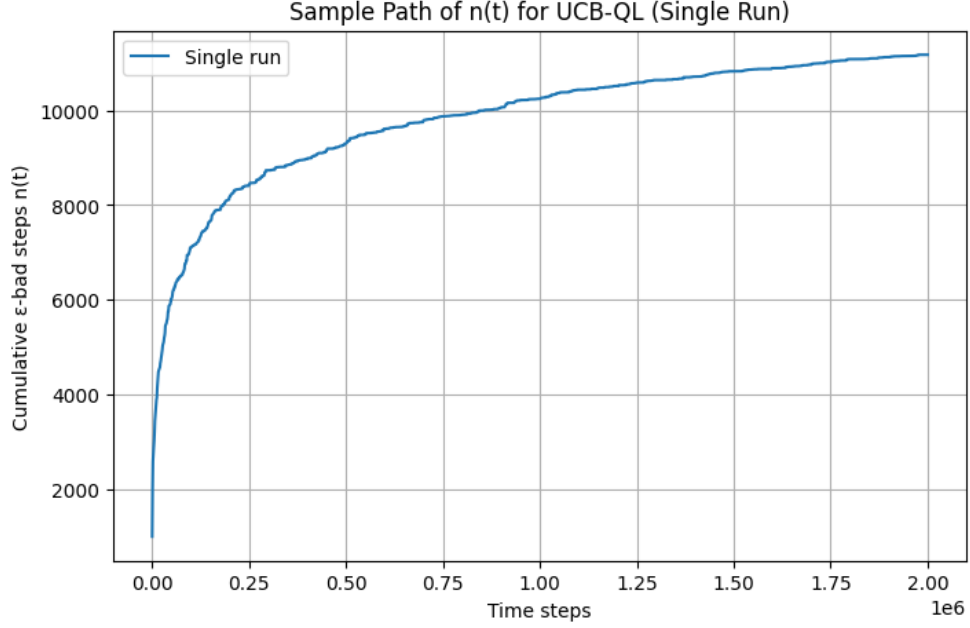


Figure 1: Sample path of the cumulative number of ε -bad steps $n(t)$ for a single run.

(ii) Averaged $n(t)$ over 100 Runs

In addition, the algorithm was executed over 100 independent runs. Figure 2 shows the average $n(t)$ with 95% confidence intervals. Notably, the individual runs exhibit very similar behavior, as evidenced by the tight confidence bounds.

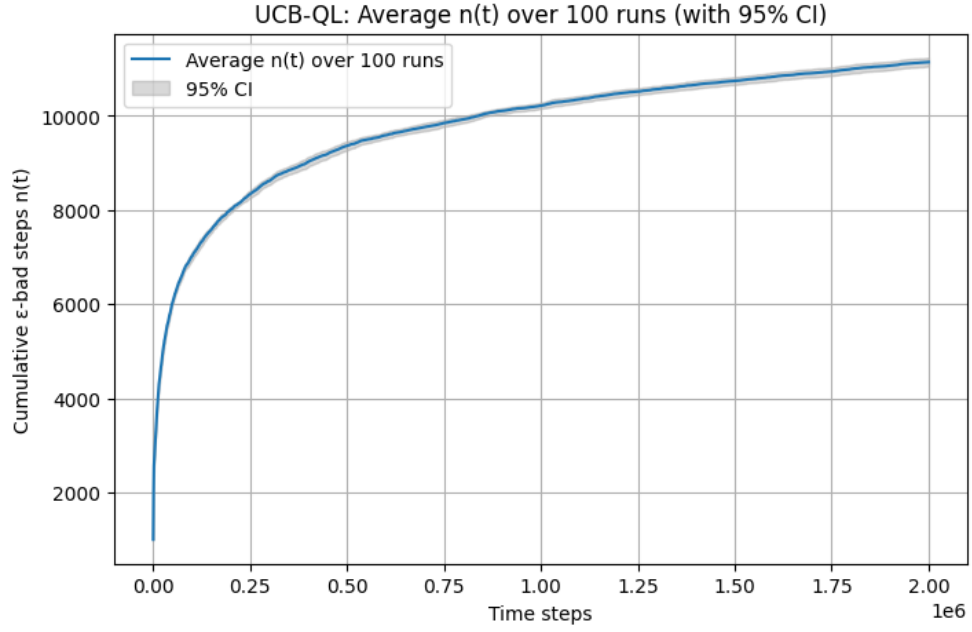


Figure 2: Averaged $n(t)$ over 100 runs with 95% confidence intervals.

In conclusion, the experimental results confirm that UCB Q-learning consistently maintains a similar behavior across runs, with the cumulative ε -bad steps growing as expected in the RiverSwim environment.