Online and Reinforcement Learning (2025)
Home Assignment 7

Davide Marchi 777881

Contents

1 Short Questions

1. **True.** *Justification:* In a finite average-reward MDP with a finite diameter the MDP is communicating (in fact, ergodic if every state is reachable from every other under some policy). This implies that the long-run average reward (gain) is independent of the starting state.

2. **False.** *Justification:* Finite diameter guarantees that there exists a policy which can reach any state from any other in a finite expected number of steps. However, this does not mean that every arbitrary choice of actions will eventually reach every state.

3. **False.** *Justification:* In an ergodic MDP the optimal gain is unique (i.e., independent of the initial state), but the optimal bias function is determined only up to an additive constant. Hence, the bias is not uniquely defined in an absolute sense.

4. **False.** *Justification:* A PAC-MDP algorithm guarantees that the number of ε-bad (i.e., non-ε-optimal) steps is bounded with high probability. However, it does not imply that there exists a finite time after which every subsequent policy is ε-optimal; occasional exploration may still yield suboptimal actions.

2 Offline Evaluation of Bandit Algorithms - the Practical Part

3 Grid-World: Continual and undiscounted

In this question, we model the 4-room grid-world as an average-reward MDP and solve it using Value Iteration (VI).

(i) Implementation of Value Iteration for Average-Reward MDPs

The following Python function implements VI for average-reward MDPs. The algorithm computes an approximation of the optimal gain g^*, the bias function b^* (normalized such that $\min_s b^*(s) = 0$), and the optimal policy.

```
def average_reward_vi(mdp, epsilon=1e-6, max_iter=10000):

    # Initialize the value function arbitrarily (here, zeros)
    V = np.zeros(mdp.nS)

    for iteration in range(max_iter):
        V_next = np.zeros(mdp.nS)
        # Update V_next for each state s using the Bellman operator:
        # V_next(s) = max_a [ R(s, a) + sum_x P(x|s,a) * V(x) ]
        for s in range(mdp.nS):
            action_values = []
```

```
            for a in range(mdp.nA):
                q_sa = mdp.R[s, a] + np.dot(mdp.P[s, a], V)
                action_values.append(q_sa)
            V_next[s] = max(action_values)

        # Compute the difference (increment) vector
        diff = V_next - V
        # The span (max difference minus min difference) is our stopping
criterion
        span_diff = np.max(diff) - np.min(diff)
        V = V_next.copy()

        if span_diff < epsilon:
            print(f"Converged after {iteration+1} iterations with span {
span_diff:.2e}.")
            break
    else:
        print("Warning: Maximum iterations reached without full
convergence.")

    # Estimate the optimal gain g* as the average of the maximum and
minimum differences.
    gain = 0.5 * (np.max(diff) + np.min(diff))
    # The bias function is approximated by normalizing V (bias is defined
    up to an additive constant)
    bias = V - np.min(V)

    # Derive the optimal policy: for each state, choose the action
maximizing:
    # Q(s, a) = R(s, a) + sum_x P(x|s,a) * V(x)
    policy = np.zeros(mdp.nS, dtype=int)
    for s in range(mdp.nS):
        best_val = -np.inf
        best_a = 0
        for a in range(mdp.nA):
            q_sa = mdp.R[s, a] + np.dot(mdp.P[s, a], V)
            if q_sa > best_val:
                best_val = q_sa
                best_a = a
        policy[s] = best_a

    return policy, gain, bias
```

Listing 1: Average-Reward Value Iteration Function

(ii) Visualization of the Optimal Policy

Running the above function on the grid-world produced the following output:

- **Optimal Gain:** $g^* = 0.07563$

- **Span of the Optimal Bias:** $\mathrm{sp}(b^*) = 0.92437$

The optimal policy for the grid-world is represented by the following table. The arrows indicate the action chosen in each state (with walls marked as "Wall"):

Wall	Wall	Wall	Wall	Wall	Wall	Wall
Wall	→	→	→	↓	↓	Wall
Wall	↓	↑	Wall	↓	←	Wall
Wall	↓	Wall	Wall	↓	Wall	Wall
Wall	→	→	→	→	↓	Wall
Wall	→	↑	Wall	→	↑	Wall
Wall	Wall	Wall	Wall	Wall	Wall	Wall

Table 1: Optimal Policy Grid: Arrows indicate the optimal actions.

(iii) Interpretation of $1/g^*$

In this grid-world, the optimal gain g^* represents the long-run average reward per time step. Hence, the quantity

$$\frac{1}{g^*} \approx \frac{1}{0.07563} \approx 13.22$$

can be interpreted as the average number of steps needed to collect one unit of reward under the optimal policy. In other words, the agent gathers one reward approximately every 13 time steps, reflecting the efficiency of the optimal strategy in this continual setting.

4 An Empirical Evaluation of UCB Q-learning